

TissueNet: Detect Lesions in Cervical Biopsies


Kirill Brodt

Task description


- Image classification task
- ~1k train images with extremely high resolution
 - ~150,000 x 85,000 pixels (~30 Gb RAM)
- 4 ordinal classes
- Weighted Class Score
- <https://www.drivendata.org/competitions/67/competition-cervical-biopsy/>

Very big images

C07_B016_S21 (80,128 x 87,296)



C07_B089_S21 (80,384 x 71,424)



Ordinal classes

- 0: benign (normal or subnormal)
- 1: low malignant potential (low grade squamous intraepithelial lesion)
- 2: high malignant potential (high grade squamous intraepithelial lesion)
- 3: invasive cancer (invasive squamous carcinoma)

Performance metric

ERROR TABLE


	Class 0 (pred)	Class 1 (pred)	Class 2 (pred)	Class 3 (pred)
Class 0 (actual)	0.0	0.1	0.7	1.0
Class 1 (actual)	0.1	0.0	0.3	0.7
Class 2 (actual)	0.7	0.3	0.0	0.3
Class 3 (actual)	1.0	0.7	0.3	0.0

Image downsampling


- It's not possible (at least now) to work with such big images which do not fit ordinary computer with 16-32 or even 64 Gb RAM
- We need to downsample it with minimal loss of the information
- Fixed 16x downsampled resolution ($\sim 5,000 \times 5,000$ --- $10,000 \times 10,000$)

Tiles grid


C07_B016_S21, page=p4 (5008x5456 pixels)



grid with 128x128 tiles



stacked top 144 tiles (1536x1536 pixels)



Tiles grid

- 36 tiles with 256 x 256 size
- 64 tiles with 192 x 192 size
- 144 tiles with 128 x 128 size

In all cases we have image with 1,536 x 1,536 input size


Issue: Huge image with white pixels

Insane huge sample

C13_B054_S11

with a lot of white pixels

294,144 x 272,128




Issue: different downsampling levels

16x (18,384 x 17,008)




8x (36,768 x 34,016)




Issue: different downsampling levels

4x (73,536 x 68,032)



2x (147,072 x 136,064)



Classes encoding




- 0 -> [0, 0, 0]: benign (normal or subnormal)
- 1 -> [1, 0, 0]: low malignant potential (low grade squamous intraepithelial lesion)
- 2 -> [1, 1, 0]: high malignant potential (high grade squamous intraepithelial lesion)
- 3 -> [1, 1, 1]: invasive cancer (invasive squamous carcinoma)

Neural Network Training

- Mixed precision
- EfficientNet-B0
- Binary Cross Entropy Loss
- Batch size 8 (20 Gb VRAM)
- AdamW with learning rate $1e-3$ or $3e-4$
- CosineAnnealing scheduler
- Augmentations on tile and whole image levels: horizontal and vertical flips, rotate on 90
- Model ensembling (mean predictions of models from 8 folds)
- <https://github.com/kbrodt/competition-cervical-biopsy>

Learning curve

~24H Nvidia V100 32GB



Further research

- Use segmentation model to extract ROIs (region of interest)
- Use adaptive downsampling based on initial image resolution